

# Final Project

M5364–Data Mining   Fall 2016   Tarleton State Univ   Dr. Scott Cook   Assigned 2016-09-30   Due 2016-12-10

This data set is called the abalone data set. Abalone's are a form of sea snails whose age can be determined by boring into the shell and, through a "boring and time consuming" process, by counting the rings much like a tree. In order to save our conchologist friends time, we will attempt to predict the number of rings in the shell by measuring:

- 1) Sex
- 2) Length
- 3) Diameter
- 4) Height
- 5) Whole Weight
- 6) Shucked Weight
- 7) Viscera weight
- 8) Shell weight

Preprocessing:

```
data = read.csv("C:/Users/Rowe/Documents/My Repo/Final Project DM/abalone.data.txt",
  header = FALSE)
# We have enough data from the UCI website to know that none of the data
# points are missing values or Nan's.
table(data[, 1])
```

```

  F      I      M
1307 1342 1528
```

```
# Because there is almost even number of 'M', 'F' and 'I' values here, there
# is minimal chance of a class imbalance problem coming from this variable.
my_shapiro(data[, 2:8])
```

```
[1] 7.442090e-29 1.648335e-28 1.181265e-47 1.013778e-27 9.340986e-32
[6] 1.777103e-29 1.565014e-28
```

```
# From this function, we can confidently say that none of the continuous
# variables are normal.
```

Beginning:

```
# This function calculates a 1x2 matrix whose values are the Root Mean
# Square Error (RMSE) and the Mean Absolute Error (MAE) respectively.
RMSE_MAE = function(true, predicted) {
  error_matrix = matrix(nrow = 1, ncol = 2)
  error_matrix[1, 1] = sqrt((1/length(true)) * sum((predicted - true)^2))
  error_matrix[1, 2] = (1/length(true)) * sum(abs(true - predicted))
  return(error_matrix)
}
# I am creating master train and test set so we can compare the accuracies
# of the models directly
set.seed(5364)
split_data = splitdata(data, 0.7)
train = split_data$train
test = split_data$test
train_rows = split_data$train_rows
```

48 Artificial Neural Networks:

```
49 Error_matrix = matrix(ncol = 2, nrow = 15)
50 for (i in 1:15) {
51   crude_ann = nnet(V9 ~ ., data = train, size = as.numeric(i), maxit = 1000,
52     linout = TRUE, trace = FALSE)
53   predicted_ann = predict(crude_ann, newdata = test)
54   Error_matrix[i, ] = RMSE_MAE(test[, 9], predicted_ann)
55 }
56 Error_matrix[argmin(Error_matrix[, 1]), ]
```

```
57 [1] 2.087510 1.464726
```

```
60 Error_matrix[argmin(Error_matrix[, 2]), ]
```

```
61 [1] 2.087510 1.464726
```

```
64 RMSE_size = argmin(Error_matrix[, 1])
65 MAE_size = argmin(Error_matrix[, 2])
66 ann_cv = function(size, d, m, data_frame) {
67   error_matrix = matrix(nrow = m + 1, ncol = 2)
68   for (i in 1:m) {
69     d_records = sample(nrow(data_frame), d)
70     temp_train = data_frame[d_records, ]
71     temp_test = data_frame[-d_records, ]
72     model = nnet(V9 ~ ., data = temp_train, size = size, maxit = 1000, linout = TRUE,
73       trace = FALSE)
74     predicted = predict(model, temp_test)
75     error_matrix[i, ] = RMSE_MAE(predicted, temp_test[, 9])
76   }
77   error_matrix[m + 1, 1] = mean(error_matrix[1:m, 1])
78   error_matrix[m + 1, 2] = mean(error_matrix[1:m, 2])
79   return(error_matrix)
80 }
81 ann_cv(RMSE_size, round(0.7 * nrow(data)), 25, data)
```

```
82      [,1]      [,2]
83 [1,] 2.095105 1.485442
84 [2,] 2.207351 1.541343
85 [3,] 2.039443 1.449004
86 [4,] 2.065539 1.448992
87 [5,] 2.064905 1.466190
88 [6,] 2.002281 1.447325
89 [7,] 2.123396 1.480538
90 [8,] 2.070027 1.465895
91 [9,] 2.086959 1.464638
92 [10,] 2.070136 1.474192
93 [11,] 2.005938 1.410567
94 [12,] 2.150864 1.485463
95 [13,] 2.177452 1.523892
96 [14,] 2.172037 1.501955
97 [15,] 1.966703 1.418677
98 [16,] 2.111945 1.465941
99 [17,] 1.962255 1.428519
100 [18,] 2.083109 1.427756
101 [19,] 1.980216 1.427262
102 [20,] 2.082509 1.491201
103 [21,] 2.059137 1.480417
104 [22,] 2.068889 1.454277
105 [23,] 2.017523 1.434391
```

```

107 [24,] 2.146891 1.493057
108 [25,] 2.134369 1.493402
109 [26,] 2.077799 1.466413
110

```

```

111 ann_cv(MAE_size, round(0.7 * nrow(data)), 25, data)

```

```

112
113      [,1]      [,2]
114 [1,]  2.829343 1.497518
115 [2,] 165.295623 9.595944
116 [3,]  2.155146 1.506562
117 [4,]  2.038264 1.465907
118 [5,]  2.036226 1.422680
119 [6,]  2.074686 1.468382
120 [7,]  2.136545 1.486556
121 [8,]  2.132773 1.493238
122 [9,]  2.097541 1.504436
123 [10,] 2.083242 1.471270
124 [11,] 2.155716 1.520520
125 [12,] 2.111108 1.493397
126 [13,] 2.071499 1.478480
127 [14,] 2.046546 1.438408
128 [15,] 2.038515 1.423895
129 [16,] 1.974339 1.396308
130 [17,] 2.105372 1.502165
131 [18,] 2.082981 1.478904
132 [19,] 2.127992 1.514478
133 [20,] 2.071086 1.455224
134 [21,] 2.243506 1.582867
135 [22,] 2.017569 1.465865
136 [23,] 2.208430 1.496113
137 [24,] 2.159128 1.542066
138 [25,] 2.192894 1.563071
139 [26,] 8.659443 1.810570
140

```

```

141 ann_n_fold_cross = function(n, data_frame, size) {
142   chopping_point = floor(nrow(data_frame)/n)
143   stopping_points = c(1/chopping_point, 1:n) * chopping_point
144   error_matrix = matrix(ncol = 2, nrow = n + 1)
145   for (i in 1:n) {
146     lower_bound = stopping_points[i]
147     upper_bound = stopping_points[i + 1]
148     temp_test = data_frame[lower_bound:upper_bound, ]
149     temp_train = data_frame[-(lower_bound:upper_bound), ]
150     model = nnet(V9 ~ ., data = temp_train, size = size, maxit = 1000, linout = TRUE,
151               trace = FALSE)
152     predicted = predict(model, temp_test)
153     error_matrix[i, ] = RMSE_MAE(predicted, temp_test[, 9])
154   }
155   error_matrix[n + 1, 1] = mean(error_matrix[1:n, 1])
156   error_matrix[n + 1, 2] = mean(error_matrix[1:n, 2])
157   return(error_matrix)
158 }
159 ann_n_fold_cross(10, data, RMSE_size)

```

```

160
161      [,1]      [,2]
162 [1,] 2.648594 1.955683
163 [2,] 3.260283 2.327803
164 [3,] 1.612628 1.217336
165 [4,] 2.050177 1.595570

```

```

166 [5,] 1.587534 1.187856
167 [6,] 2.845384 2.046782
168 [7,] 1.385622 1.054785
169 [8,] 2.188472 1.620859
170 [9,] 1.639786 1.196326
171 [10,] 1.957278 1.408840
172 [11,] 2.117576 1.561184

```

```

174 ann_n_fold_cross(10, data, MAE_size)

```

```

175
176      [,1]      [,2]
177 [1,] 2.587194 1.911639
178 [2,] 3.179929 2.274739
179 [3,] 1.722674 1.349745
180 [4,] 2.050171 1.595562
181 [5,] 1.583929 1.190128
182 [6,] 2.856398 2.044390
183 [7,] 1.349745 1.027931
184 [8,] 3.225724 2.373436
185 [9,] 1.656324 1.219490
186 [10,] 1.959487 1.417368
187 [11,] 2.217157 1.640443

```

189 Decision Trees:

```

190 # First, we're going to create a decision tree to find a good guess as how
191 # accurate we can get it.
192 first_tree = rpart(V9 ~ ., data = train)
193 first_predicted = predict(first_tree, newdata = test)
194 RMSE_MAE(first_predicted, test[, 9])

```

```

195
196      [,1]      [,2]
197 [1,] 2.452406 1.754041

```

```

199 # Now that we have an estimate for accuracy, we're going to use randomForest
200 # to get a more accurate model using the same test and train data as the
201 # decision tree did. Note: like all tune functions, it does include a
202 # 10-fold cross validation test of accuracy.
203 forestmodel = tune.randomForest(V9 ~ ., data = train)
204 forestmodel

```

```

205
206
207 Error estimation of 'randomForest' using 10-fold cross validation: 4.709902

```

```

209 randomforest_cv = function(d, m, data_frame) {
210   error_matrix = matrix(nrow = m + 1, ncol = 2)
211   for (i in 1:m) {
212     d_records = sample(nrow(data_frame), d)
213     temp_train = data_frame[d_records, ]
214     temp_test = data_frame[-d_records, ]
215     model = randomForest(V9 ~ ., data = temp_train)
216     predicted = predict(model, temp_test)
217     error_matrix[i, ] = RMSE_MAE(predicted, temp_test[, 9])
218   }
219   error_matrix[m + 1, 1] = mean(error_matrix[1:m, 1])
220   error_matrix[m + 1, 2] = mean(error_matrix[1:m, 2])
221   return(error_matrix)
222 }
223 randomforest_cv(round(0.7 * nrow(data)), 25, data)

```

```

224      [,1]      [,2]
225
226 [1,] 2.189761 1.559031
227 [2,] 2.032656 1.451119
228 [3,] 2.147769 1.502647
229 [4,] 2.160276 1.517783
230 [5,] 2.155134 1.504519
231 [6,] 2.237676 1.545699
232 [7,] 2.156094 1.521302
233 [8,] 2.065127 1.496686
234 [9,] 2.080927 1.502607
235 [10,] 2.097192 1.522242
236 [11,] 2.146814 1.506259
237 [12,] 2.029483 1.470464
238 [13,] 2.089718 1.517656
239 [14,] 2.149859 1.547441
240 [15,] 2.225827 1.560536
241 [16,] 2.151263 1.543244
242 [17,] 2.139978 1.500482
243 [18,] 2.183016 1.509353
244 [19,] 2.049982 1.459659
245 [20,] 2.161172 1.491655
246 [21,] 2.316180 1.586989
247 [22,] 2.089543 1.470437
248 [23,] 1.925467 1.424065
249 [24,] 2.140897 1.521362
250 [25,] 2.250510 1.603325
251 [26,] 2.134893 1.513462
252

```

253 Support Vector Machines:

```

254 # Here I'm making a for loop that runs the tune function five times to find
255 # a good guess at the best parameters.
256 best_parameters = matrix(ncol = 2, nrow = 5)
257 obj <- tune.svm(V9 ~ ., data = train, gamma = 2^(0:4), cost = 2^(0:4))
258 best_parameters[1, ] = as.matrix(obj$best.parameters)
259 for (i in 1:4) {
260   gamma = as.numeric(obj$best.parameters)[1]
261   cost = as.numeric(obj$best.parameters)[2]
262   obj <- tune.svm(V9 ~ ., data = train, gamma = (0.25 * gamma):(1.75 * gamma),
263     cost = (0.5 * cost):(1.5 * cost))
264   best_parameters[i + 1, ] = as.matrix(obj$best.parameters)
265 }
266 best_model = obj$best.model
267 svm_cv = function(d, m, data_frame) {
268   error_matrix = matrix(nrow = m + 1, ncol = 2)
269   for (i in 1:m) {
270     d_records = sample(nrow(data_frame), d)
271     temp_train = data_frame[d_records, ]
272     temp_test = data_frame[-d_records, ]
273     model = svm(V9 ~ ., data = temp_train, gamma = best_parameters[5, 1],
274       cost = best_parameters[5, 2])
275     predicted = as.numeric(predict(model, temp_test))
276     error_matrix[i, ] = RMSE_MAE(predicted, temp_test[, 9])
277   }
278   error_matrix[m + 1, 1] = mean(error_matrix[1:m, 1])
279   error_matrix[m + 1, 2] = mean(error_matrix[1:m, 2])
280   return(error_matrix)
281 }
282 svm_cv(round(0.7 * nrow(data)), 25, data)

```

```

283      [,1]      [,2]
284
285 [1,] 2.290255 1.554808
286 [2,] 2.192552 1.505933
287 [3,] 2.186716 1.510708
288 [4,] 2.179284 1.517026
289 [5,] 2.182027 1.498752
290 [6,] 2.219791 1.533720
291 [7,] 2.357776 1.592112
292 [8,] 2.323458 1.560145
293 [9,] 2.254139 1.545554
294 [10,] 2.300648 1.561159
295 [11,] 2.265284 1.539249
296 [12,] 2.299996 1.529469
297 [13,] 2.180669 1.470896
298 [14,] 2.240785 1.510838
299 [15,] 2.263901 1.561225
300 [16,] 2.268788 1.529638
301 [17,] 2.271419 1.541745
302 [18,] 2.218457 1.482707
303 [19,] 2.240139 1.545508
304 [20,] 2.333338 1.609474
305 [21,] 2.234419 1.524602
306 [22,] 2.338490 1.618267
307 [23,] 2.205766 1.516662
308 [24,] 2.337352 1.556476
309 [25,] 2.135988 1.479137
310 [26,] 2.252857 1.535832
311

```

```

312 obj
313
314

```

```

315 Parameter tuning of 'svm':
316
317 - sampling method: 10-fold cross validation
318
319 - best parameters:
320      gamma      cost
321 0.00390625 1.9375
322
323 - best performance: 5.145983
324

```

```

325 Naive Bayes:

```

```

326 nrings = data[, 9]
327 rings = as.factor(nrings)
328 sex = data[, 1]
329 nb_data = data[, -1]
330 nb_train = train[, -1]
331 nb_test = test[, -1]
332 Error_matrix = matrix(nrow = 14, ncol = 2)
333 # Since the Shapiro-Wilks test suggests that all of the numeric variables
334 # are nonnormal, I have to find out the best way to turn them into discrete
335 # factors using a for loop. Thus, I am keeping the train and test set
336 # constant and only changing number of levels for the newly-factorized
337 # variables.
338 colnames(Error_matrix) = c("RMSE", "MAE")
339 for (i in 2:15) {
340     discretized_train = discretizer(nb_train[, -8], i)
341     discretized_train = cbind(sex[train_rows], discretized_train)

```

```

342 discretized_test = discretizer(nb_test[, -8], i)
343 discretized_test = cbind(sex[-train_rows], discretized_test)
344 model = naiveBayes(rings[train_rows] ~ ., data = discretized_train)
345 predicted_naive_Bayes = predict(model, test)
346 predicted_naive_Bayes = as.numeric(predicted_naive_Bayes)
347 Error_matrix[i - 1, ] = RMSE_MAE(predicted_naive_Bayes, test[, 9])
348 }
349 Error_matrix[argmin(Error_matrix[, 1]), ]

```

RMSE	MAE
3.547534	2.709497

```

354 Error_matrix[argmin(Error_matrix[, 2]), ]

```

RMSE	MAE
3.662323	2.604948

```

359 # We can see that the factor levels that minimizes the Root Mean Square
360 # error and the Mean Absolute error are very close but different and both of
361 # them don't change all that much when switching between the two. So, we
362 # could confidently pick either one and would receive similar results. I am
363 # now going to use 10-fold cross-validation to double check my numbers.
364 data_RMSE = discretizer(nb_data[, -8], argmin(Error_matrix[, 1]))
365 data_RMSE = cbind(sex, data_RMSE)
366 data_MAE = discretizer(nb_data[, -8], argmin(Error_matrix[, 2]))
367 data_MAE = cbind(sex, data_MAE)
368 data = cbind(sex, data)
369 naiveBayes_cv = function(d, m, data_frame) {
370   error_matrix = matrix(nrow = m + 1, ncol = 2)
371   for (i in 1:m) {
372     d_records = sample(nrow(data_frame), d)
373     temp_test = data_frame[d_records, ]
374     temp_train = data_frame[-d_records, ]
375     model = naiveBayes(rings[-d_records] ~ ., data = temp_train)
376     predicted = as.numeric(predict(model, temp_test))
377     error_matrix[i, ] = RMSE_MAE(predicted, nrings[d_records])
378   }
379   error_matrix[m + 1, 1] = mean(error_matrix[1:m, 1])
380   error_matrix[m + 1, 2] = mean(error_matrix[1:m, 2])
381   return(error_matrix)
382 }
383 naiveBayes_cv(round(0.7 * nrow(data_RMSE)), 25, data_RMSE)

```

	[,1]	[,2]
[1,]	3.060773	2.007182
[2,]	3.401939	2.136115
[3,]	3.450897	2.155609
[4,]	3.254482	2.045828
[5,]	3.449410	2.161081
[6,]	3.033838	1.975718
[7,]	3.504297	2.223324
[8,]	3.248224	2.117305
[9,]	3.370275	2.136457
[10,]	3.421436	2.158345
[11,]	2.901837	1.914501
[12,]	3.318583	2.152531
[13,]	3.087362	2.022914
[14,]	2.921160	1.927839
[15,]	3.483693	2.117647

```

401 [16,] 3.427777 2.121067
402 [17,] 3.236569 2.064295
403 [18,] 2.952253 1.929207
404 [19,] 3.490411 2.189124
405 [20,] 3.168868 2.058824
406 [21,] 3.285336 2.110807
407 [22,] 3.155403 2.049590
408 [23,] 2.829696 1.879959
409 [24,] 3.365401 2.092681
410 [25,] 3.086531 2.023940
411 [26,] 3.236258 2.070876

```

```

413 naiveBayes_cv(round(0.7 * nrow(data_MAE)), 25, data_MAE)

```

```

414
415      [,1]      [,2]
416 [1,] 3.011322 1.985978
417 [2,] 3.144491 2.037620
418 [3,] 3.029495 2.009576
419 [4,] 3.050924 1.981190
420 [5,] 3.057866 2.020862
421 [6,] 3.150847 2.102257
422 [7,] 3.083594 2.044118
423 [8,] 3.015464 1.987004
424 [9,] 3.121732 2.037278
425 [10,] 3.077377 2.019494
426 [11,] 3.058985 1.993502
427 [12,] 3.053669 2.015048
428 [13,] 3.073818 2.011286
429 [14,] 3.190701 2.067031
430 [15,] 3.012514 1.978796
431 [16,] 3.085866 2.027360
432 [17,] 3.190433 2.049590
433 [18,] 3.079376 2.021546
434 [19,] 3.086254 2.037278
435 [20,] 3.029100 2.003762
436 [21,] 3.043461 2.017100
437 [22,] 2.996521 1.964090
438 [23,] 3.256793 2.127223
439 [24,] 3.266492 2.122777
440 [25,] 3.044416 2.005130
441 [26,] 3.088461 2.026676

```

```

443 nb_n_fold_cross = function(n, data_frame) {
444   chopping_point = floor(nrow(data_frame)/n)
445   stopping_points = c(1/chopping_point, 1:n) * chopping_point
446   error_matrix = matrix(ncol = 2, nrow = n + 1)
447   for (i in 1:n) {
448     lower_bound = stopping_points[i]
449     upper_bound = stopping_points[i + 1]
450     temp_test = data_frame[lower_bound:upper_bound, ]
451     temp_train = data_frame[-(lower_bound:upper_bound), ]
452     model = naiveBayes(rings[-(lower_bound:upper_bound)] ~ ., data = temp_train)
453     predicted = as.numeric(predict(model, temp_test))
454     error_matrix[i, ] = RMSE_MAE(predicted, nrings[lower_bound:upper_bound])
455   }
456   error_matrix[n + 1, 1] = mean(error_matrix[1:n, 1])
457   error_matrix[n + 1, 2] = mean(error_matrix[1:n, 2])
458   return(error_matrix)
459 }

```



```
nb_n_fold_cross(10, data_RMSE)
```

```
      [,1]      [,2]
[1,] 4.236120 3.074519
[2,] 5.044064 3.930622
[3,] 3.055311 1.602871
[4,] 2.486807 1.270335
[5,] 2.845293 1.507177
[6,] 4.653033 3.516746
[7,] 2.730739 1.337321
[8,] 3.635945 2.497608
[9,] 2.833920 1.653110
[10,] 3.049041 1.827751
[11,] 3.457027 2.221806
```

```
nb_n_fold_cross(10, data_MAE)
```

```
      [,1]      [,2]
[1,] 3.964688 2.843750
[2,] 5.021961 3.861244
[3,] 1.599192 1.093301
[4,] 1.537412 1.143541
[5,] 1.547494 1.078947
[6,] 4.469460 3.401914
[7,] 1.409978 1.069378
[8,] 3.138357 2.174641
[9,] 2.336009 1.476077
[10,] 2.664672 1.775120
[11,] 2.768922 1.991791
```

Knn function:

```
# Because knn can't calculate distance with factor variables, we have to
# create two dummy variables. Also, because knn relies on distance, I scaled
# the data so that variables who have smaller ranges aren't more important.
male = as.numeric(data[, 1] == "M")
female = as.numeric(data[, 1] == "F")
knn_data = cbind(male, female, data[, -1])
knn_train = cbind(male[train_rows], female[train_rows], train[, -1])
knn_test = cbind(male[-train_rows], female[-train_rows], test[, -1])
Error_matrix = matrix(ncol = 2, nrow = 50)
for (i in 1:50) {
  predicted_knn = as.numeric(knn(knn_train[, -10], knn_test[, -10], knn_train[,
    10], k = i))
  Error_matrix[i, ] = cbind(RMSE_MAE(knn_test[, 10], predicted_knn))
}
Error_matrix[argmin(Error_matrix[, 1]), ]
```

```
[1] 3.519414 2.777334
```

```
Error_matrix[argmin(Error_matrix[, 2]), ]
```

```
[1] 3.585686 2.774142
```

```
k_RMSE = argmin(Error_matrix[, 1])
k_MAE = argmin(Error_matrix[, 2])
# Delete-d cross validation
knn_cv = function(d, m, data_frame, k) {
  error_matrix = matrix(nrow = m + 1, ncol = 2)
```

```

518   for (i in 1:m) {
519     d_records = sample(nrow(data_frame), d)
520     temp_test = data_frame[-d_records, ]
521     temp_train = data_frame[d_records, ]
522     predicted_knn = as.numeric(knn(temp_train[, -10], temp_test[, -10],
523       cl = temp_train[, 10], k = k))
524     error_matrix[i, ] = RMSE_MAE(temp_test[, 10], predicted_knn)
525   }
526   error_matrix[m + 1, 1] = mean(error_matrix[1:m, 1])
527   error_matrix[m + 1, 2] = mean(error_matrix[1:m, 2])
528   return(error_matrix)
529 }
530 knn_cv(round(0.7 * nrow(knn_data)), 25, knn_data, k = k_RMSE)

```

```

531 Warning in knn(temp_train[, -10], temp_test[, -10], cl = temp_train[, 10], : NAs introduced by coercion
532

```

```

533 Warning in knn(temp_train[, -10], temp_test[, -10], cl = temp_train[, 10], : NAs introduced by coercion
534

```

```

535 Error in knn(temp_train[, -10], temp_test[, -10], cl = temp_train[, 10], : NA/NaN/Inf in foreign function
536 call (arg 6)

```

```

541 knn_cv(round(0.7 * nrow(knn_data)), 25, knn_data, k = k_MAE)

```

```

542 Warning in knn(temp_train[, -10], temp_test[, -10], cl = temp_train[, 10], : NAs introduced by coercion
543

```

```

544 Warning in knn(temp_train[, -10], temp_test[, -10], cl = temp_train[, 10], : NAs introduced by coercion
545

```

```

546 Error in knn(temp_train[, -10], temp_test[, -10], cl = temp_train[, 10], : NA/NaN/Inf in foreign function
547 call (arg 6)

```

```

552 # 10-fold cross validation

```

```

553 knn_n_fold_cross = function(n, data_frame, k) {
554   chopping_point = floor(nrow(data_frame)/n)
555   stopping_points = c(1/chopping_point, 1:n) * chopping_point
556   error_matrix = matrix(ncol = 2, nrow = n + 1)
557   for (i in 1:n) {
558     lower_bound = stopping_points[i]
559     upper_bound = stopping_points[i + 1]
560     temp_test = data_frame[lower_bound:upper_bound, ]
561     temp_train = data_frame[-(lower_bound:upper_bound), ]
562     predicted = as.numeric(knn(temp_train[, -10], temp_test[, -10], cl = temp_train[,
563       10], k = k))
564     error_matrix[i, ] = RMSE_MAE(predicted, temp_test[, 10])
565   }
566   error_matrix[n + 1, 1] = mean(error_matrix[1:n, 1])
567   error_matrix[n + 1, 2] = mean(error_matrix[1:n, 2])
568   return(error_matrix)
569 }
570 knn_n_fold_cross(10, knn_data, k_RMSE)

```

```

571 Warning in knn(temp_train[, -10], temp_test[, -10], cl = temp_train[, 10], : NAs introduced by coercion
572

```

```

573 Warning in knn(temp_train[, -10], temp_test[, -10], cl = temp_train[, 10], : NAs introduced by coercion
574

```

```

575 Error in knn(temp_train[, -10], temp_test[, -10], cl = temp_train[, 10], : NA/NaN/Inf in foreign function
576 call (arg 6)

```

```

581 knn_n_fold_cross(10, knn_data, k_MAE)

```

```
Warning in knn(temp_train[, -10], temp_test[, -10], cl = temp_train[, 10], : NAs introduced by coercion
```

```
Warning in knn(temp_train[, -10], temp_test[, -10], cl = temp_train[, 10], : NAs introduced by coercion
```

```
Error in knn(temp_train[, -10], temp_test[, -10], cl = temp_train[, 10], : NA/NaN/Inf in foreign function
call (arg 6)
```

Weighted Knn-function:

```
k_vector = c(1:50)
# Since the optimal kernel does not always provide the best result, all of
# them are being tested with k values between 1 and 50. Once the optimal
# kernel and k value are calculated, they will cross validated.
optimal_test = train.kknn(V9 ~ ., data = knn_data, ks = k_vector, kernel = "optimal",
  ykernel = -1.2:1.2)
triangular_test = train.kknn(V9 ~ ., data = knn_data, ks = k_vector, kernel = "triangular",
  ykernel = -1.2:1.2)
gaussian_test = train.kknn(V9 ~ ., data = knn_data, ks = k_vector, kernel = "gaussian",
  ykernel = -1.2:1.2)
epanechnikov_test = train.kknn(V9 ~ ., data = knn_data, ks = k_vector, kernel = "epanechnikov",
  ykernel = -1.2:1.2)
biweight_test = train.kknn(V9 ~ ., data = knn_data, ks = k_vector, kernel = "biweight",
  ykernel = -1.2:1.2)
triweight_test = train.kknn(V9 ~ ., data = knn_data, ks = k_vector, kernel = "triweight",
  ykernel = -1.2:1.2)
cosine_test = train.kknn(V9 ~ ., data = knn_data, ks = k_vector, kernel = "cos",
  ykernel = -1.2:1.2)
inverted_test = train.kknn(V9 ~ ., data = knn_data, ks = k_vector, kernel = "inv",
  ykernel = -1.2:1.2)
rectangular_test = train.kknn(V9 ~ ., data = knn_data, ks = k_vector, kernel = "rectangular",
  ykernel = -1.2:1.2)
optimal_data_frame = rbind(c("optimal", min(optimal_test$MEAN.ABS), min(optimal_test$MEAN.SQU),
  optimal_test$best.parameters$k), c("triangular", min(triangular_test$MEAN.ABS),
  min(triangular_test$MEAN.SQU), triangular_test$best.parameters$k), c("epanechnikov",
  min(epanechnikov_test$MEAN.ABS), min(epanechnikov_test$MEAN.SQU), epanechnikov_test$best.parameters$k),
  c("biweight", min(biweight_test$MEAN.ABS), min(biweight_test$MEAN.SQU),
  biweight_test$best.parameters$k), c("triweight", min(triweight_test$MEAN.ABS),
  min(triweight_test$MEAN.SQU), triweight_test$best.parameters$k), c("cos",
  min(cosine_test$MEAN.ABS), min(cosine_test$MEAN.SQU), cosine_test$best.parameters$k),
  c("inv", min(inverted_test$MEAN.ABS), min(inverted_test$MEAN.SQU), inverted_test$best.parameters$k),
  c("gaussian", min(gaussian_test$MEAN.ABS), min(gaussian_test$MEAN.SQU),
  gaussian_test$best.parameters$k), c("rectangular", min(rectangular_test$MEAN.ABS),
  min(rectangular_test$MEAN.SQU), rectangular_test$best.parameters$k))
# Since train.kknn includes a leave-one-out cross validation technique, I
# will use that method.
optimal_data_frame[argmin(optimal_data_frame[, 2]), ]
```

```
[1] "gaussian"          "1.51924144935422" "4.8439507372699"
[4] "18"
```

```
optimal_data_frame[argmin(optimal_data_frame[, 3]), ]
```

```
[1] "biweight"          "1.52023104377889" "4.82416270611321"
[4] "43"
```

```
optimal_k_abs = optimal_data_frame[argmin(optimal_data_frame[, 2]), 4]
optimal_k_abs = as.numeric(optimal_k_abs)
optimal_kernel_abs = optimal_data_frame[argmin(optimal_data_frame[, 2]), 1]
```

```

642 optimal_kernel_abs = as.character(optimal_kernel_abs)
643 optimal_k_rmse = optimal_data_frame[argmin(optimal_data_frame[, 3]), 4]
644 optimal_k_rmse = as.numeric(optimal_k_rmse)
645 optimal_kernel_rmse = optimal_data_frame[argmin(optimal_data_frame[, 3]), 1]
646 optimal_kernel_rmse = as.character(optimal_kernel_rmse)
647 # 10 fold cross validation using convenience command
648 fold_RMSE = cv.kknn(V9 ~ ., data = knn_data, kcv = 10, k = optimal_k_rmse, kernel = optimal_kernel_rmse)
649 RMSE_MAE(fold_RMSE[[1]][1:4177], fold_RMSE[[1]][4178:(2 * 4177)])

```

```

650           [,1]      [,2]
651 [1,]  2.206253  1.526847
652
653

```

```

654 fold_MAE = cv.kknn(V9 ~ ., data = knn_data, kcv = 10, k = optimal_k_abs, kernel = optimal_kernel_abs)
655 RMSE_MAE(fold_MAE[[1]][1:4177], fold_MAE[[1]][4178:(2 * 4177)])

```

```

656           [,1]      [,2]
657 [1,]  2.205438  1.528683
658
659

```